





FORNAX

# Introduction to Deep Learning, part 2

- Rafał Cycoń, FORNAX


[WWW.FORNAX.CO](http://WWW.FORNAX.CO)

# Image classification





# Localization






# Photorealistic style transfer



# Image generation



(a) User constraints  $v_g$  at different update steps



$G(z_0)$

(b) Updated images according to user edits



$G(z_1)$




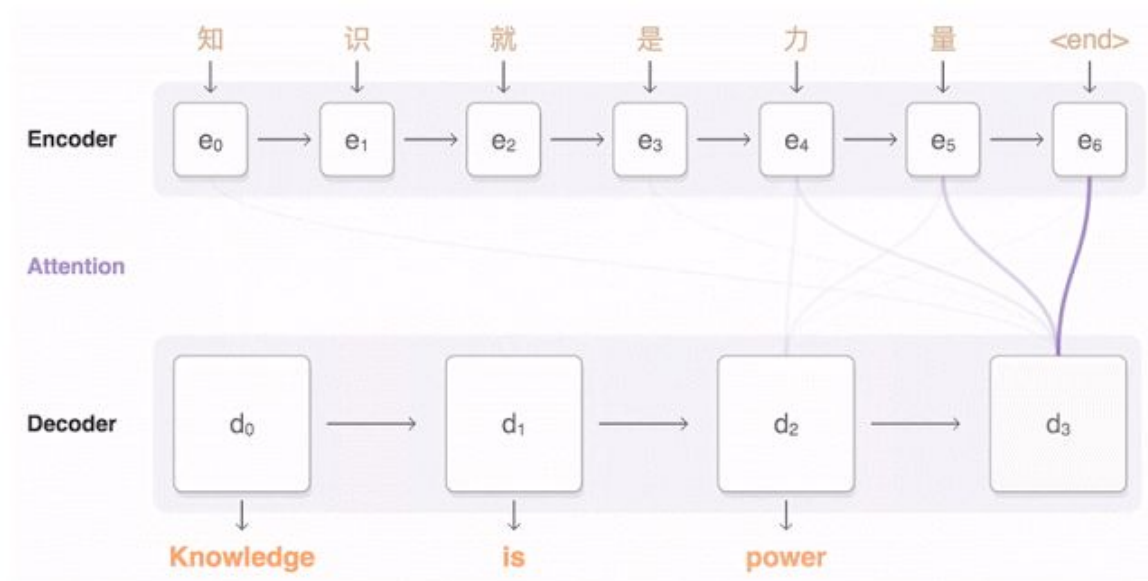
(c) Linear interpolation between  $G(z_0)$  and  $G(z_1)$



# Sound synthesis



# Translation



# Text generation

---

PANDARUS:

Alas, I think he shall be come approached and the day  
When little strain would be attain'd into being never fed,  
And who is but a chain and subjects of his death,  
I should not sleep.

Second Senator:

They are away this miseries, produced upon my soul,  
Breaking and strongly should be buried, when I perish  
The earth and thoughts of many states.

DUKE VINCENTIO:

Well, your wit is in the care of side and that.

Second Lord:

They would be ruled after this chamber, and  
my fair nudes begun out of the fact, to be conveyed,  
Whose noble souls I'll have the heart of the wars.

Clown:

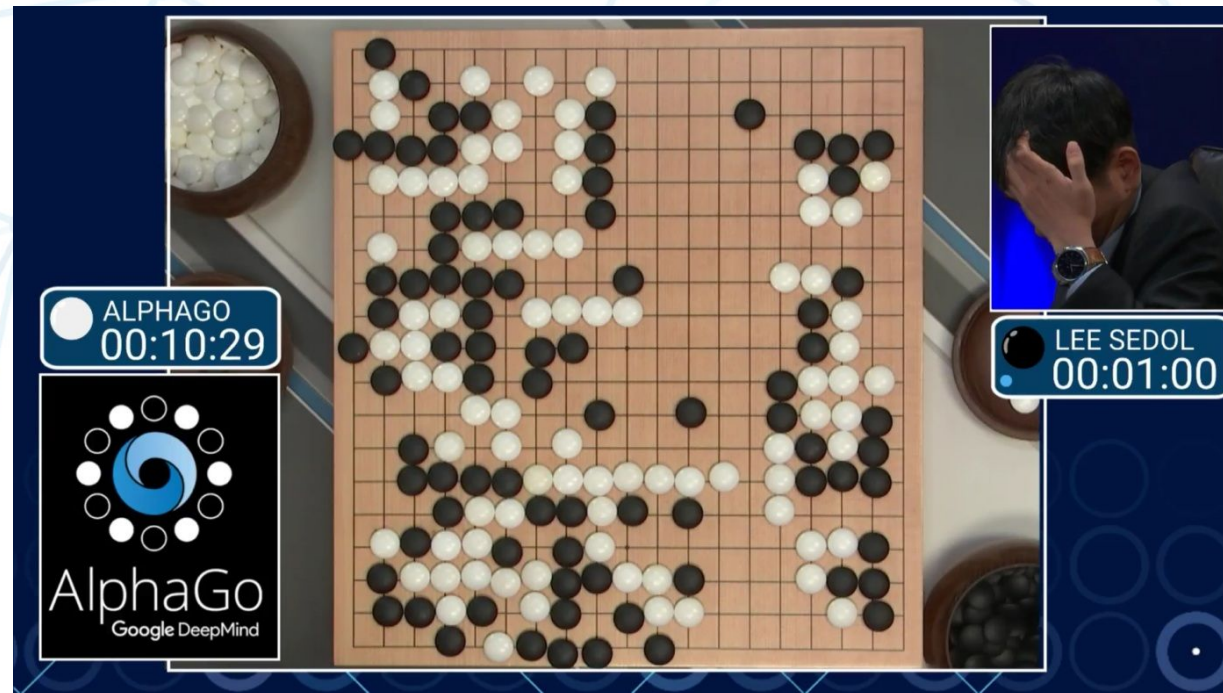
Come, sir, I will make did behold your worship.

VIOLA:

I'll drink it.




# Reinforcement learning



# ImageNet



ILSVRC top-5 error on ImageNet





# Challenges in image processing

Viewpoint variation



Scale variation



Deformation



Occlusion



Illumination conditions



Background clutter




Intra-class variation



# Convolution

$$\begin{aligned}(f * g)(t) &\stackrel{\text{def}}{=} \int_{-\infty}^{\infty} f(\tau) g(t - \tau) d\tau \\ &= \int_{-\infty}^{\infty} f(t - \tau) g(\tau) d\tau.\end{aligned}$$

$$\begin{aligned}(f * g)[n] &\stackrel{\text{def}}{=} \sum_{m=-\infty}^{\infty} f[m] g[n - m] \\ &= \sum_{m=-\infty}^{\infty} f[n - m] g[m].\end{aligned}$$



35	40	41	45	50
40	40	42	46	52
42	46	50	55	55
48	52	56	58	60
56	60	65	70	75

$\times$

	0	1	0	
	0	0	0	
	0	0	0	

$=$

		42		



# Convolution

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

1	0	1
0	1	0
1	0	1


1 <sub>x1</sub>	1 <sub>x0</sub>	1 <sub>x1</sub>	0	0
0 <sub>x0</sub>	1 <sub>x1</sub>	1 <sub>x0</sub>	1	0
0 <sub>x1</sub>	0 <sub>x0</sub>	1 <sub>x1</sub>	1	1
0	0	1	1	0
0	1	1	0	0

Image


4		

Convolved  
Feature

# Convolution




1	1	1
1	1	1
1	1	1




0	1	0
1	-4	1
0	1	0



# Convolution




# Convolution







# Convolution Layer




6 filters, each of size  $5 \times 5 \times 3$

Convolution Layer



# ConvNet










object models



object parts  
(combination  
of edges)




edges




pixels



# Pooling



# Pooling



Discards 75% of info!




# Convolution

---



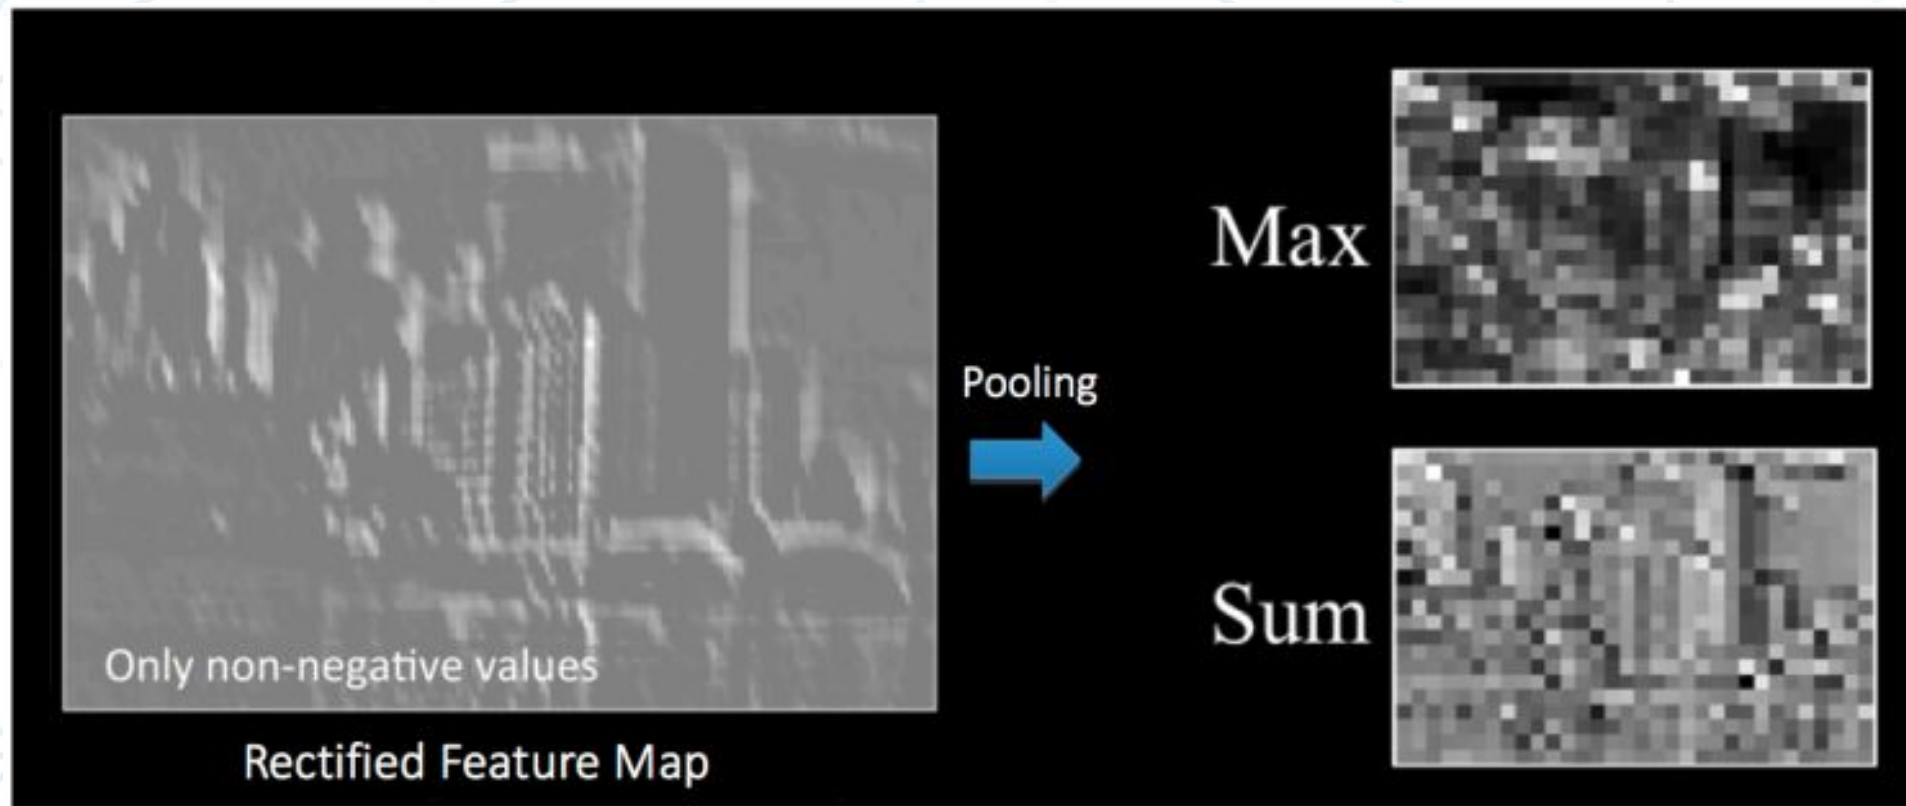
Input

# Nonlinearity






# Pooling



# AlexNet






# VGG-16

ConvNet Configuration					
A	A-LRN	B	C	D	E
11 weight layers	11 weight layers	13 weight layers	16 weight layers	16 weight layers	19 weight layers
input ( $224 \times 224$ RGB image)					
conv3-64	conv3-64 <b>LRN</b>	conv3-64 <b>conv3-64</b>	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64
maxpool					
conv3-128	conv3-128	conv3-128 <b>conv3-128</b>	conv3-128 conv3-128	conv3-128 conv3-128	conv3-128 conv3-128
maxpool					
conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256 <b>conv1-256</b>	conv3-256 conv3-256 <b>conv3-256</b>	conv3-256 conv3-256 conv3-256 <b>conv3-256</b>
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 <b>conv1-512</b>	conv3-512 conv3-512 <b>conv3-512</b>	conv3-512 conv3-512 conv3-512 <b>conv3-512</b>
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 <b>conv1-512</b>	conv3-512 conv3-512 <b>conv3-512</b>	conv3-512 conv3-512 conv3-512 <b>conv3-512</b>
maxpool					
FC-4096					
FC-4096					
FC-1000					
soft-max					

# Smaller filters

---





# Smaller filters – what do we gain?

Consider an input with height  $H$ , width  $W$ ,  $C$  channels.  
Assume stride=1, padding such that  $H$  and  $W$  do not change on the output.

A single 7x7 filter:

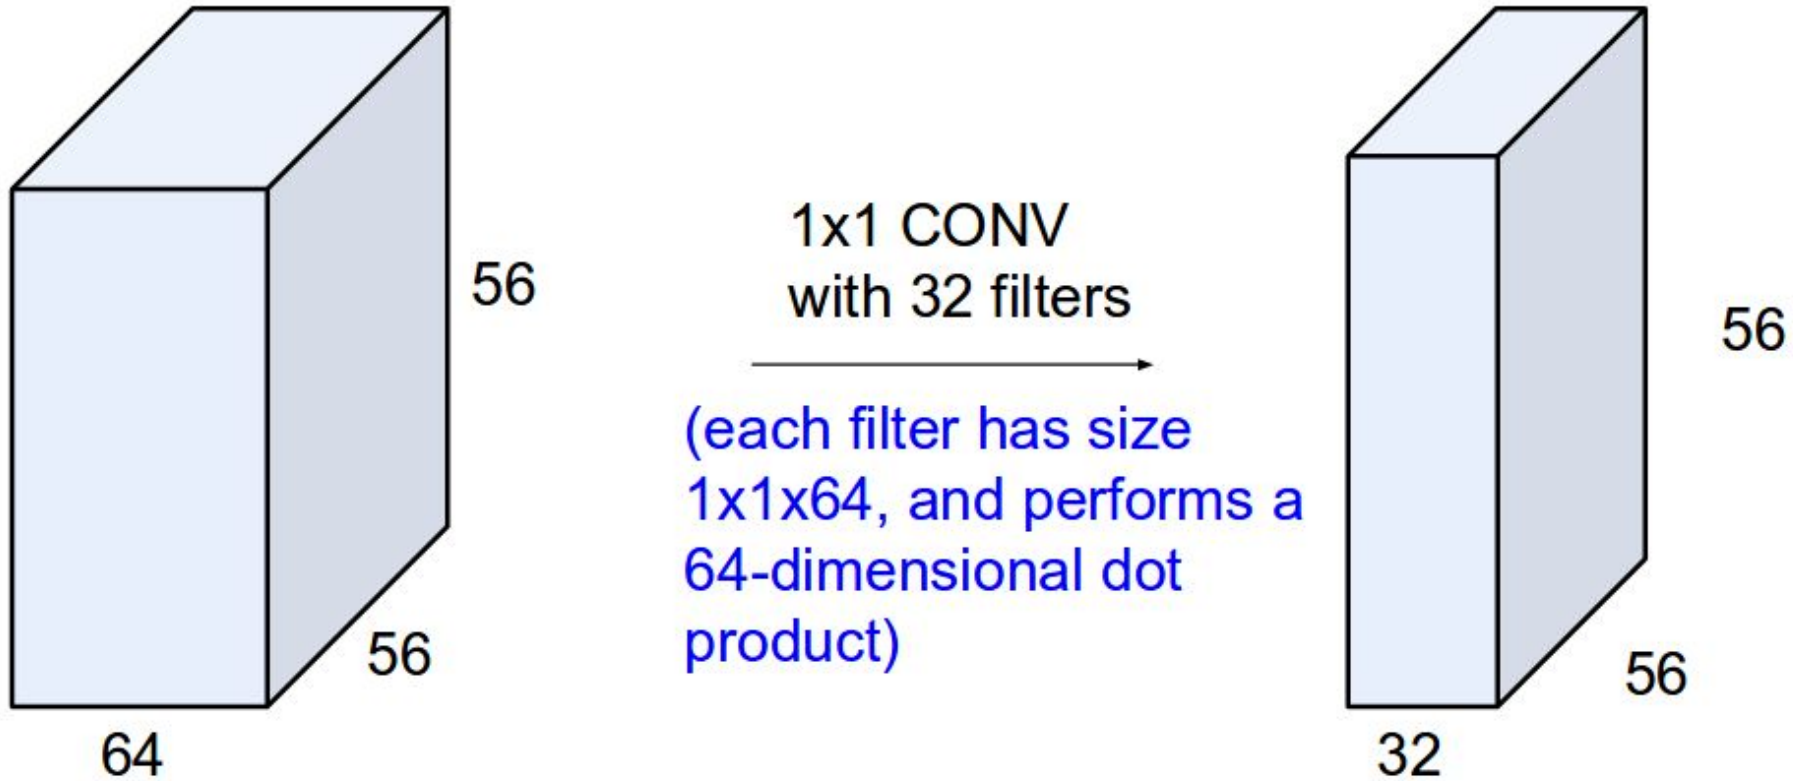
- parameters:  $C \cdot (7 \cdot 7 \cdot C) = 49C^2$  (and a bias)
- operations:  $(H \cdot W \cdot C) \cdot (7 \cdot 7 \cdot C) = 49 HWC^2$

A stack of three 3x3 filters (same receptive field):

- parameters:  $3 \cdot C \cdot (3 \cdot 3 \cdot C) = 27C^2$  (and a bias)
- operations:  $(H \cdot W \cdot C) \cdot 3 \cdot (3 \cdot 3 \cdot C) = 27 HWC^2$

- Less parameters
- Less computations
- More nonlinearity

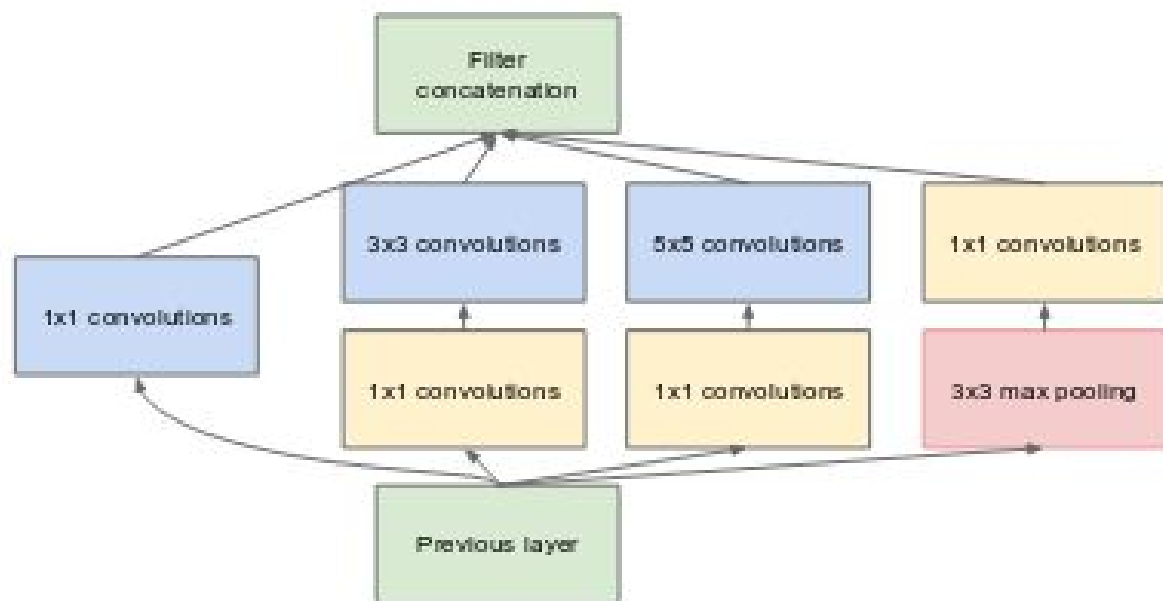
# 1x1 filters



Even less params, less computations, more nonlinearity  
Can be used to increase/decrease number of channels (bottlenecks)

# GoogLeNet

<https://arxiv.org/abs/1409.4842>



(b) Inception module with dimensionality reduction






# Global pooling

---

First appeared in “Network in Network” (<https://arxiv.org/pdf/1312.4400.pdf>)

Facilitates the use of all conv networks (ConvNet can work with all input resolutions!)



# ResNet

<https://arxiv.org/abs/1512.03385>






Figure 1. Training error (left) and test error (right) on CIFAR-10 with 20-layer and 56-layer “plain” networks. The deeper network has higher training error, and thus test error. Similar phenomena on ImageNet is presented in Fig. 4.



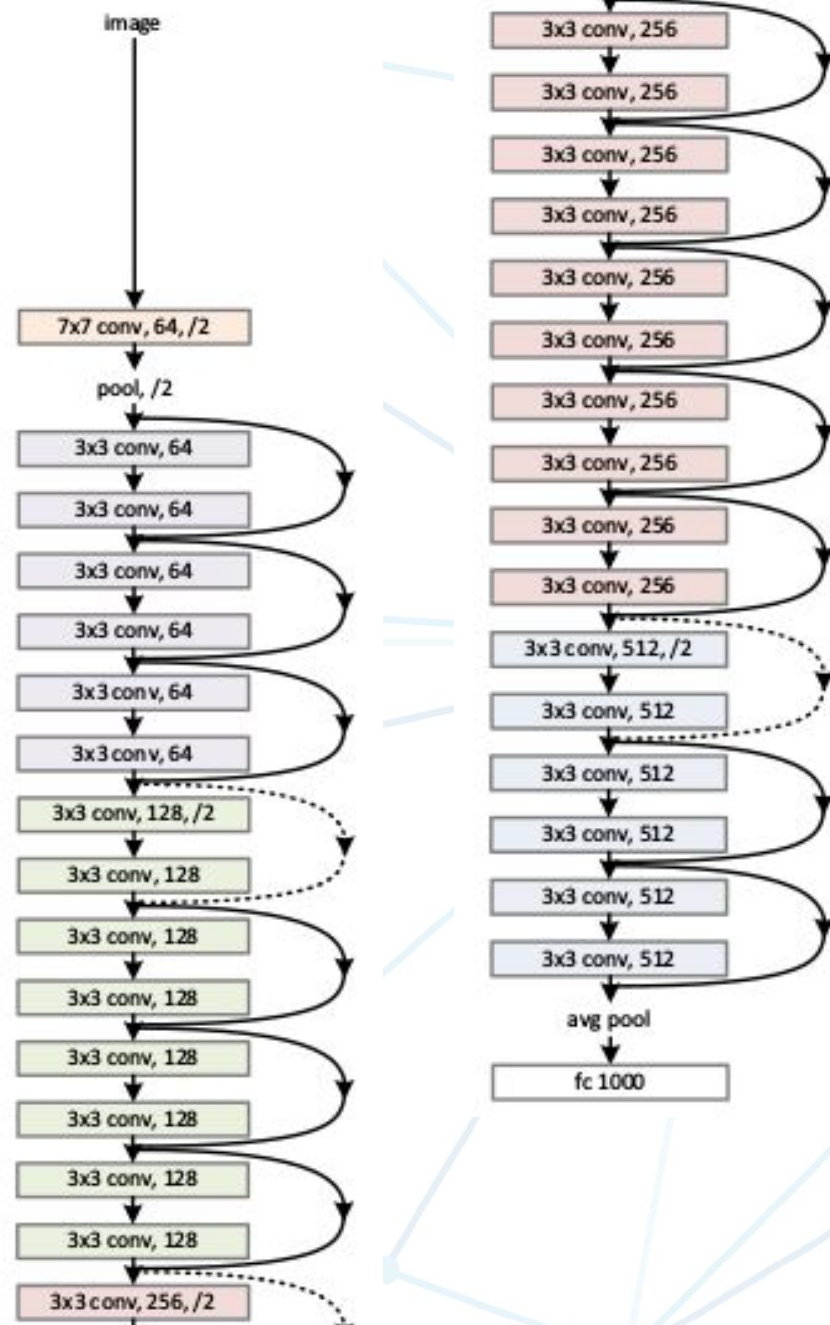
# ResNet

<https://arxiv.org/abs/1512.03385>



“Bottleneck” building block



## 34-layer residual





# ResNet v2

<https://arxiv.org/abs/1603.05027>



# WideResNet

<https://arxiv.org/abs/1605.07146>

	depth-k	# params	CIFAR-10	CIFAR-100
NIN [20]			8.81	35.67
DSN [19]			8.22	34.57
FitNet [24]			8.39	35.04
Highway [28]			7.72	32.39
ELU [5]			6.55	24.28
original-ResNet[11]	110	1.7M	6.43	25.16
	1202	10.2M	7.93	27.82
stoc-depth[14]	110	1.7M	5.23	24.58
	1202	10.2M	4.91	-
pre-act-ResNet[13]	110	1.7M	6.37	-
	164	1.7M	5.46	24.33
	1001	10.2M	4.92(4.64)	22.71
WRN (ours)	40-4	8.9M	4.53	21.18
	16-8	11.0M	4.27	20.43
	28-10	36.5M	<b>4.00</b>	<b>19.25</b>

Model	top-1 err, %	top-5 err, %	#params	time/batch 16
ResNet-50	24.01	7.02	25.6M	49
ResNet-101	22.44	6.21	44.5M	82
ResNet-152	22.16	6.16	60.2M	115
<b>WRN-50-2-bottleneck</b>	21.9	6.03	68.9M	93
pre-ResNet-200	21.66	5.79	64.7M	154

# Simple tips

---

**ConvNets need lots of data**

**Use data augmentation or transfer learning**

**Lots of parameters = overfitting**

**Regularize! L2, dropout, BatchNorm**

**Use residual blocks and bottleneck convs (e.g.  $1 \times 1 \rightarrow 3 \times 3 \rightarrow 1 \times 1$ )**

**Global pooling prior to softmax (or any other dense layer)**

**Gives a possibility to process images of arbitrary resolution**

**We recommend using Adam as a state-of-the-art optimization algorithm**



# Transfer learning

---

Get a model trained on a similar dataset.

Got “lots” of data?




**Yes!**

Fine-tune last few layers, freeze the rest

**No...**

Train a (simple) classifier on network output and/or on features (i.e. filter maps) from different layers

# Debugging





# FORNAX

**Rafał Cycoń**  
Co-founder, CTO

+48 607 697 054  
[rafal.cycon@fornax.co](mailto:rafal.cycon@fornax.co)

[WWW.FORNAX.CO](http://WWW.FORNAX.CO)